

PCA_Mini_Project_Solution

January 28, 2026

0.0.1 PCA Mini Project - Solution

In the lesson, you saw how you could use PCA to substantially reduce the dimensionality of the handwritten digits. In this mini-project, you will be using the `cars.csv` file.

To begin, run the cell below to read in the necessary libraries and the dataset. I also read in the helper functions that you used throughout the lesson in case you might find them helpful in completing this project. Otherwise, you can always create functions of your own!

```
[1]: import pandas as pd
import numpy as np
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score
from helper_functions import do_pca, scree_plot, plot_components, pca_results
from IPython import display
import test_code2 as t

import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline

df = pd.read_csv('./data/cars.csv')
```

1. Now your data is stored in `df`. Use the below cells to take a look your dataset. At the end of your exploration, use your findings to match the appropriate variable to each key in the dictionary below.

```
[2]: df.head()
```

```
[2]:
```

	Sports	SUV	Wagon	Minivan	Pickup	AWD	RWD	\
Acura 3.5 RL	0	0	0	0	0	0	0	
Acura 3.5 RL Navigation	0	0	0	0	0	0	0	
Acura MDX	0	1	0	0	0	1	0	
Acura NSX S	1	0	0	0	0	0	1	
Acura RSX	0	0	0	0	0	0	0	

	Retail	Dealer	Engine	Cylinders	Horsepower	\
Acura 3.5 RL	43755	39014	3.5	6	225	
Acura 3.5 RL Navigation	46100	41100	3.5	6	225	
Acura MDX	36945	33337	3.5	6	265	
Acura NSX S	89765	79978	3.2	6	290	
Acura RSX	23820	21761	2.0	4	200	
	CityMPG	HighwayMPG	Weight	Wheelbase	Length	Width
Acura 3.5 RL	18	24	3880	115	197	72
Acura 3.5 RL Navigation	18	24	3893	115	197	72
Acura MDX	17	23	4451	106	189	77
Acura NSX S	17	24	3153	100	174	71
Acura RSX	24	31	2778	101	172	68

[3] : df.describe()

	Sports	SUV	Wagon	Minivan	Pickup	AWD	\
count	387.000000	387.000000	387.000000	387.000000	387.0	387.000000	
mean	0.116279	0.152455	0.072351	0.054264	0.0	0.201550	
std	0.320974	0.359926	0.259404	0.226830	0.0	0.401677	
min	0.000000	0.000000	0.000000	0.000000	0.0	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.0	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.0	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	0.0	0.000000	
max	1.000000	1.000000	1.000000	1.000000	0.0	1.000000	
	RWD	Retail	Dealer	Engine	Cylinders		\
count	387.000000	387.000000	387.000000	387.000000	387.000000		
mean	0.242894	33231.180879	30440.653747	3.127390	5.757106		
std	0.429387	19724.634576	17901.179282	1.014314	1.490182		
min	0.000000	10280.000000	9875.000000	1.400000	3.000000		
25%	0.000000	20997.000000	19575.000000	2.300000	4.000000		
50%	0.000000	28495.000000	26155.000000	3.000000	6.000000		
75%	0.000000	39552.500000	36124.000000	3.800000	6.000000		
max	1.000000	192465.000000	173560.000000	6.000000	12.000000		
	Horsepower	CityMPG	HighwayMPG	Weight	Wheelbase		\
count	387.000000	387.000000	387.000000	387.000000	387.000000		
mean	214.444444	20.312661	27.263566	3532.457364	107.211886		
std	70.262822	5.262333	5.636005	706.003622	7.086553		
min	73.000000	10.000000	12.000000	1850.000000	89.000000		
25%	165.000000	18.000000	24.000000	3107.000000	103.000000		
50%	210.000000	19.000000	27.000000	3469.000000	107.000000		
75%	250.000000	21.500000	30.000000	3922.000000	112.000000		
max	493.000000	60.000000	66.000000	6400.000000	130.000000		

```
Length      Width
count    387.000000  387.000000
mean     184.961240  71.276486
std      13.237999  3.368329
min     143.000000  64.000000
25%    177.000000  69.000000
50%    186.000000  71.000000
75%    193.000000  73.000000
max     221.000000  81.000000
```

```
[4]: df.shape
```

```
[4]: (387, 18)
```

```
[5]: a = 7
b = 66
c = 387
d = 18
e = 0.23
f = 0.05
```

```
solution_1_dict = {
    'The number of cars in the dataset': c,
    'The number of car features in the dataset': d,
    'The number of dummy variables in the dataset': a,
    'The proportion of minivans in the dataset': f,
    'The max highway mpg for any car': b
}
```

```
[6]: # Check your solution against ours by running this cell
display.HTML(t.check_question_one(solution_1_dict))
```

Nice job! Looks like your dataset matches what we found!

```
[6]: <IPython.core.display.HTML object>
```

2. There are some particularly nice properties about PCA to keep in mind. Use the dictionary below to match the correct variable as the key to each statement. When you are ready, check your solution against ours by running the following cell.

```
[7]: a = True
b = False

solution_2_dict = {
    'The components span the directions of maximum variability.': a,
    'The components are always orthogonal to one another.': a,
    'Eigenvalues tell us the amount of information a component holds': a
```

```
}
```

```
[8]: # Check your solution against ours by running this cell  
t.check_question_two(solution_2_dict)
```

That's right these are all true. Principal components are orthogonal, span the directions of maximum variability, and the corresponding eigenvalues tell us how much of the original variability is explained by each component.

3. Fit PCA to reduce the current dimensionality of the dataset to 3 dimensions. You can use the helper functions, or perform the steps on your own. If you fit on your own, be sure to standardize your data. At the end of this process, you will want an X matrix with the reduced dimensionality to only 3 features. Additionally, you will want your **pca** object back that has been used to fit and transform your dataset.

```
[9]: pca, X_pca = do_pca(3, df)
```

4. Once you have your **pca** object, you can take a closer look at what comprises each of the principal components. Use the **pca_results** function from the **helper_functions** module assist with taking a closer look at the results of your analysis. The function takes two arguments: the full dataset and the **pca** object you created.

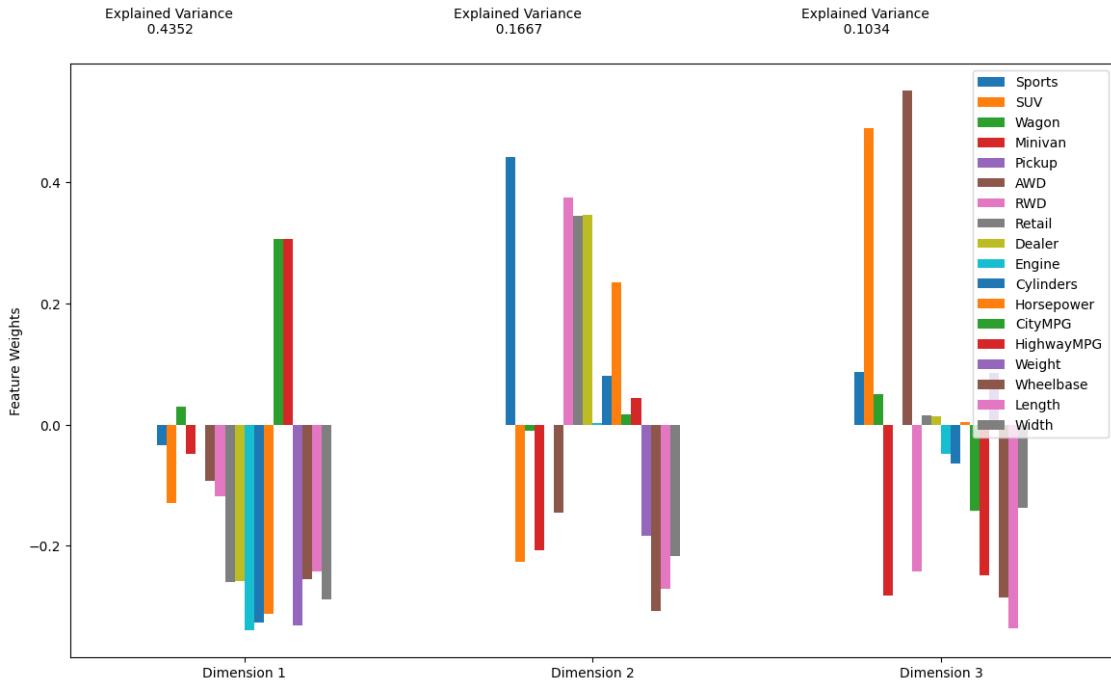
```
[10]: pca_results(df, pca)
```

```
[10]:
```

	Explained Variance	Sports	SUV	Wagon	Minivan	Pickup	\
Dimension 1	0.4352	-0.0343	-0.1298	0.0289	-0.0481	-0.0	
Dimension 2	0.1667	0.4420	-0.2261	-0.0106	-0.2074	0.0	
Dimension 3	0.1034	0.0875	0.4898	0.0496	-0.2818	0.0	

	AWD	RWD	Retail	Dealer	Engine	Cylinders	Horsepower	\
Dimension 1	-0.0928	-0.1175	-0.2592	-0.2576	-0.3396	-0.3263	-0.3118	
Dimension 2	-0.1447	0.3751	0.3447	0.3453	0.0022	0.0799	0.2342	
Dimension 3	0.5506	-0.2416	0.0154	0.0132	-0.0489	-0.0648	0.0040	

	CityMPG	HighwayMPG	Weight	Wheelbase	Length	Width	
Dimension 1	0.3063	0.3061	-0.3317	-0.2546	-0.2414	-0.2886	
Dimension 2	0.0169	0.0433	-0.1832	-0.3066	-0.2701	-0.2163	
Dimension 3	-0.1421	-0.2486	0.0851	-0.2846	-0.3361	-0.1369	



5. Use the results, to match each of the variables as the value to the most appropriate key in the dictionary below. When you are ready to check your answers, run the following cell to see if your solution matches ours!

```
[11]: a = 'car weight'
b = 'sports cars'
c = 'gas mileage'
d = 0.4352
e = 0.3061
f = 0.1667
g = 0.7053

solution_5_dict = {
    'The first component positively weights items related to': c,
    'The amount of variability explained by the first component is': d,
    'The largest weight of the second component is related to': b,
    'The total amount of variability explained by the first three components': g
}
```

```
[12]: # Run this cell to check if your solution matches ours.
t.check_question_five(solution_5_dict)
```

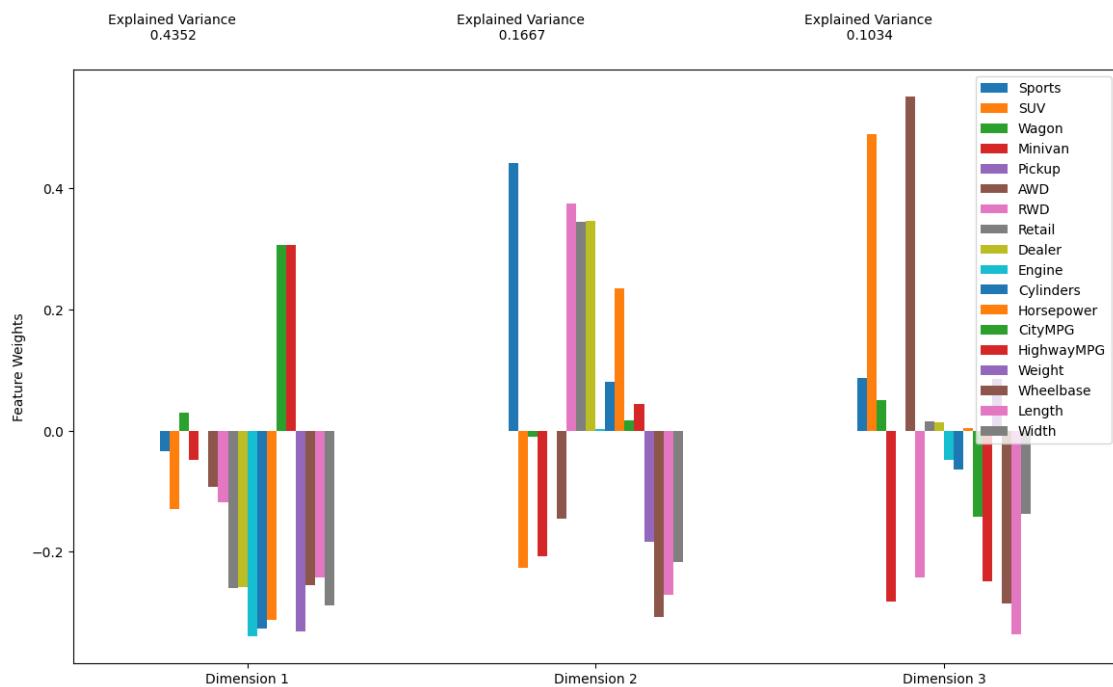
That's right! Looks like you know a lot about PCA!

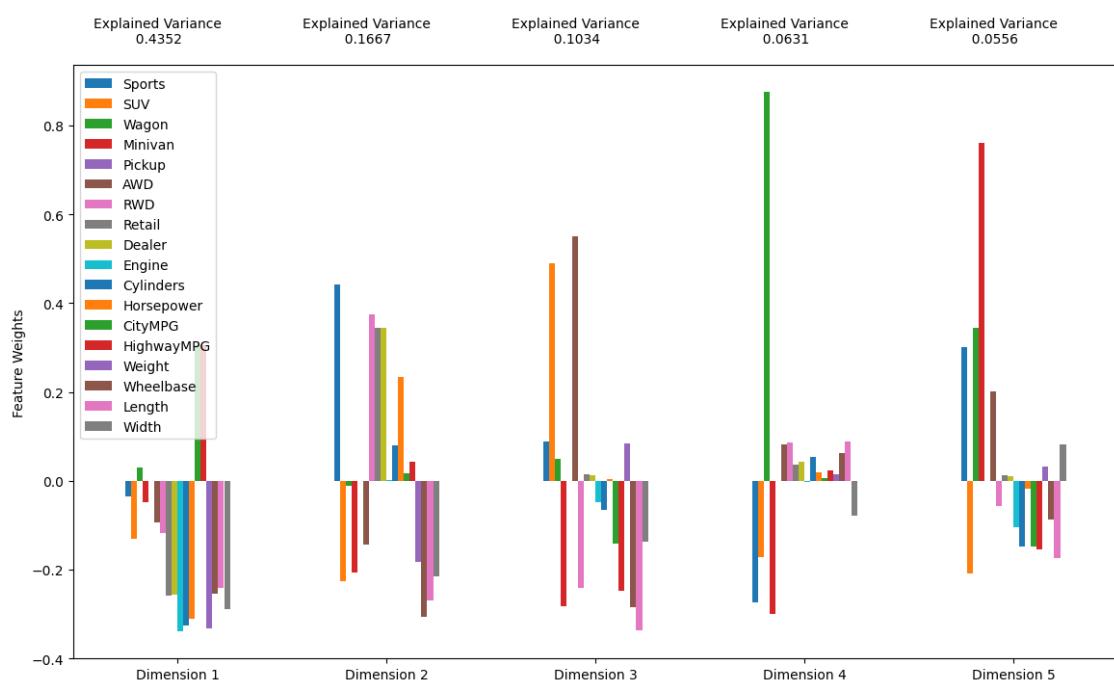
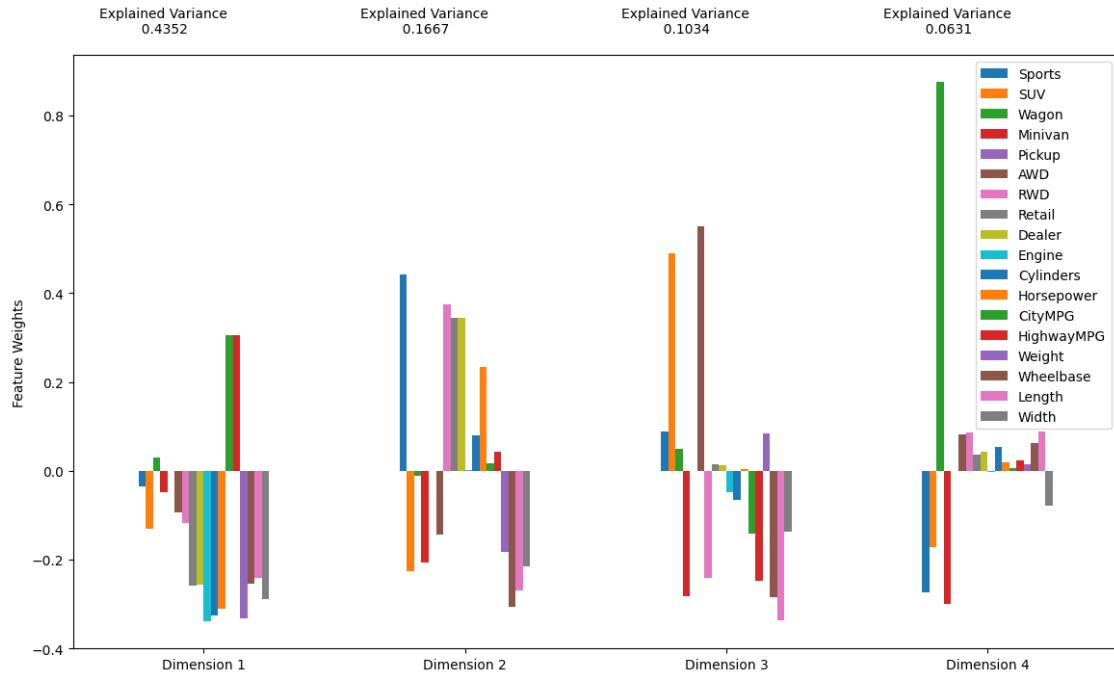
6. How many components need to be kept to explain at least 85% of the variability in the original dataset? When you think you have the answer, store it in the variable `num_comps`. Then run the

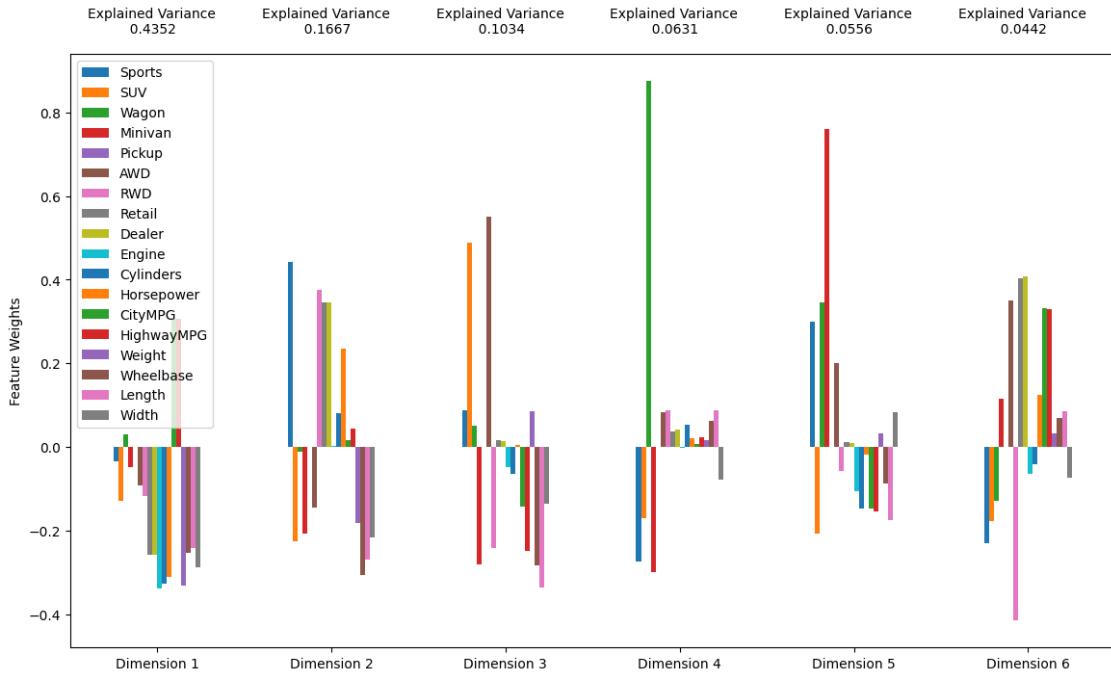
following cell to see if your solution matches ours!

```
[13]: for comp in range(3, df.shape[1]):  
    pca, X_pca = do_pca(comp, df)  
    comp_check = pca_results(df, pca)  
    if comp_check['Explained Variance'].sum() > 0.85:  
        break  
  
num_comps = comp_check.shape[0]  
print("Using {} components, we can explain {}% of the variability in the  
original data.".format(comp_check.shape[0], comp_check['Explained Variance'].  
sum()))
```

Using 6 components, we can explain 0.8682000000000001% of the variability in the original data.







```
[14]: # How check your answer here to complete this mini project
display.HTML(t.question_check_six(num_comps))
```

Nice job! That's right! With 6 components, you can explain more than 85% of the variability in the original dataset.

```
[14]: <IPython.core.display.HTML object>
```

```
[ ]:
```